

# CSE446: Blockchain & Cryptocurrencies

## Lecture – 5: Distributed Systems & Consensus Algorithms



Inspiring Excellence

# Agenda

---

- FLP Impossibility, CAP Theorem
- Various consensus algorithms

# FLP Impossibility

---


- A well-known theorem\*, by Fischer, Lynch and Paterson, called FLP Impossibility has shown that
  - a deterministic consensus protocol cannot satisfy all three properties in an asynchronous network
- Safety and liveness are favoured over fault tolerance in the domain of distributed system applications

\*M. J. Fischer, N. A. Lynch, and M. S. Paterson. "Impossibility of distributed consensus with one faulty process". Journal of the ACM (JACM), 32(2):374382, 1985

# CAP Theorem

- CAP theorem\* states that a shared replicated datastore (or, more generally, a replicated state machine) cannot achieve both consistency and availability when a network partitions in such a way that an arbitrary number of messages might be dropped

Pick any two



<https://hazelcast.com/wp-content/uploads/2021/12/cap-theorem-diagram-800x753-1.png>

\*Brewer, E. A. "Towards robust distributed systems.". In PODC (Vol. 7), July 2000.

# Crash-tolerant consensus algorithm

- Algorithms belonging to this class aim to guarantee the atomic broadcast (total order) of messages within the participating nodes in the presence of certain number of node failures
- These algorithms utilise the notion of views or epochs
  - which imply a certain duration of time
- A leader is selected for each epoch who takes decisions regarding the atomic broadcast, and all other nodes comply with its decision
- In case a leader fails due to a crash failure, the protocols elect a new leader to function

# Crash-tolerant consensus algorithm

- The best known algorithms belonging to this class can continue to function if the following condition holds:
  - $f < n/2$  where  $f$  is the number of faulty nodes and  $n$  is the total number of participating nodes
- Examples of some well-known crash-tolerant consensus protocol are
  - Paxos, Viewstamped Replication, ZooKeeper, and Raft

# RAFT

---

- Raft is a distributed consensus algorithm
- It solves the problem of getting multiple servers to agree on a shared state even in the face of failures
- The shared status is usually a data structure supported by a replicated log
- We need the system to be fully operational as long as a majority of the servers are up

# RAFT

---


- Raft works by electing a leader in the cluster
- The leader is responsible for accepting client requests and managing the replication of the log to other servers
- The data flows only in one direction: leader -> other servers
- Raft decomposes consensus into three sub-problems
  - Leader Election: A new leader needs to be elected in case of the failure of an existing one
  - Log replication: The leader needs to keep the logs of all servers in sync with its own through replication
  - Safety: to ensure consistency across all nodes




# RAFT: leader election

- An epoch is termed as a term in RAFT
- A leader is selected in each term
- If a candidate wins the election, it remains the leader for the rest of the term
  - If the vote is split, then that term ends without a leader
- The term number increases monotonically and is also exchanged in every communication

Each node can only be a leader, follower, or candidate



# RAFT: leader election



# RAFT: log replication

---

- A client communicates with the leader to update the log (append to the log)
- The leader updates its log entry and replicates the value to the followers
- Each follower updates their log entries and replies back to the leader
- The leader commits the value to the log if the majority of the nodes reply back
- The leader replicates this commitment to the followers and they also commit the value

# RAFT: safety

---

- RAFT ensures that it can provide consistency (safety) and availability when the network partitions using the following rule:
  - If one of the previous leaders has committed a log entry at a particular index, no other leader can apply a different log entry for that index
- This is assuming that this condition holds:  $f < n/2$
- A fantastic visualisation of RAFT can be viewed from this location: <https://thesecretlivesofdata.com/raft>

# Byzantine consensus algorithm

- Byzantine algorithms aim to reach consensus amid of certain nodes exhibiting Byzantine behaviour
- Such Byzantine nodes are assumed to be under the control of an adversary and behave unpredictably with malicious intent
- Similar to any crash-tolerant consensus protocol, these protocols also utilise the concept of views/epochs
- A leader is elected in each view to order messages for atomic broadcast
  - other honest nodes are assumed to follow the instructions from the leader

# Practical Byzantine Fault Tolerance (PBFT)

- One of the most well-known algorithms under this class is called Practical Byzantine Fault Tolerant (PBFT)
- It can achieve consensus in the presence of a certain number of Byzantine nodes under an eventual synchronous network assumption
- PBFT uses cryptographic algorithms such as signature, signature verification, and hash to ensure that everything stays irrevocable, unforgeable, and indisputable

# Practical Byzantine Fault Tolerance (PBFT)

- The following conditions are required
- $f < n/3$ , where  $f$  is the number of Byzantine nodes and  $n$  denotes the number of total nodes participating in the network, that is at least  $n = 3f + 1$ 
  - A distributed system that contains  $3f + 1$  nodes can have at most  $f$  faulty nodes (byzantine nodes)

# PBFT: roles

---

- An epoch is known as a view (a consensus round) in PBFT
- There are three different roles in PBFT
- Client:
  - Client nodes send transaction requests
- Primary:
  - These are main nodes which act as leaders
  - They initiate the consensus mechanism
  - In each view, there is one and only one primary node




# PBFT: roles

---

- Replica:
  - These nodes are responsible for processing each request
  - There are many replica nodes and they all act in the same fashion
- Both Primary and Replica nodes are considered as consensus nodes, because they participate in the consensus process

# PBFT: steps


- Client sends a request to the system
- Consensus nodes (Replica1, Replica2, Replica3 & Replica4) receive the request
- Then the system participates in a three-phase consensus protocol: pre-prepare, prepare and commit



<https://media.geeksforgeeks.org/wp-content/uploads/pbft-algorithm-communication-visual.png>

# PBFT: steps – pre~prepare


- Primary node verifies the request and generates a pre-prepare message which is broadcast to all replica nodes
- After receiving the messages, each replica node will validate the message
- If validated, each replica node will broadcast a corresponding prepare message



<https://media.geeksforgeeks.org/wp-content/uploads/pbft-algorithm-communication-visual.png>

# PBFT: steps – prepare


- Each node waits until  $2f+1$  prepare messages are received
- Then each node will broadcast commit messages to other nodes
  - message implying that it is ready for processing the request



<https://media.geeksforgeeks.org/wp-content/uploads/pbft-algorithm-communication-visual.png>

# PBFT: steps – commit

- Each node waits until  $2f+1$  commit messages are received
- Then, the node processes each request, which changes the system states
- Each node replies back to the client
- The primary is rotated when there is a proposal from the replicas to do so or the primary remains unavailable for a duration
  - This mechanism is called a view change



<https://media.geeksforgeeks.org/wp-content/uploads/pbft-algorithm-communication-visual.png>

